


Hard Shadows

Filtering



Shadow Map Filtering





Unfiltered



3x3 PCF

Shadow Map Filtering

- Depth Values are not “blendable”
 - Traditional bilinear filtering is inappropriate
 - Interpolating depths



3.2	1.5
-----	-----

$\text{linearFilter}(3.2, 1.5) = 2.8$


$2.8 < 3$

shadowed

Percentage Closer Filtering


[Reeves et al. 1987]

- Average comparison results, not depth values



Percentage Closer Filtering - Practice


- 2D filter kernel
- Here 3x3
- Need shadow test result before filtering
 - Pre-filtering does not work
- Depth texture uses 2x2 PCF if GL_LINEAR



Percentage Closer Filtering - Results



scene



standard shadow maps



PCF 2x2



PCF 3x3



PCF 4x4



PCF 4x4 + bilinear lookups

PCF + Bilinear Lookups

- Nearest neighbour lookups
→ quantization artifacts
- Bilinear lookups
 - For 2x2 kernel size straight forward
 - For bigger kernel sizes
 - Poisson disk + bilinear lookups (hack)
 - Exact bicubic convolution using 2x2 bilinear lookups [Hadwiger, GPU Gems 2]

Percentage Closer Filtering


- Good quality needs big kernel size
 - Bandlimiting for oversampled areas!
- Big kernel size is slow (quadratic growth)
- Pre-filtering desirable
 - Only filter once
 - Less texture lookups
 - Big kernels cheaper

Variance Shadow Maps

[Donnelly and Lauritzen 2006]

- Estimate PCF outcome with statistics
 - A representation that filters linearly
 - Use **mean** and **variance** of depth samples inside kernel
 - Can be calculated from linearized depth map


4	4	1	1	1
4	1	1	1	4
1	1	1	1	4
1	4	4	4	4



Variance Shadow Maps

[Donnelly and Lauritzen 2006]

- Estimate PCF outcome with statistics
 - Really want a cumulative distribution function (CDF) of a set of depths
 - $F(t) = P(x \leq t)$
 - $F(t)$ is the probability that a fragment at distance “ t ” from the light is in shadow



Variance Shadow Maps

[Donnelly and Lauritzen 2006]

- **Mean** = $\mu = E(x)$
- **Variance** = $\sigma^2 = E(x^2) - E(x)^2$
- Approximate fraction of distribution that is more distant than shaded point d ($P(x \geq d)$)
- Prefiltering (for a certain kernel size)
 - $E(x)$ on kernel simple to calculate from input depth texture
 - $E(x^2)$ on kernel simple to calculate from input depth texture squared
- Estimate PCF shadow test outcome with Chebyshev's Inequality

Variance Shadow Maps

[Donnelly and Lauritzen 2006]

- Chebyshev's inequality
 - Given a shadow map depth value distribution with mean and variance (*for certain kernel*), the probability $P(x \geq d)$ that a random depth value z drawn from this distribution (*this kernel*) is greater or equal to a given depth value d (*current fragment depth*) has an upper bound of

$$p(d) = \frac{\sigma^2}{\sigma^2 + (d - \mu)^2} \geq P(x \geq d)$$

- Percentage of fragments over the filter region that are more distant than the current fragment

Variance Shadow Maps


[Donnelly and Lauritzen 2006]

- Inequality only gives an upper bound
 - Becomes equality for single planar occluder and receiver
 - In small neighbourhood
occluder and receiver have constant depth and
thus $p(d)$ will provide a close approximation to $P(x \geq d)$
- In practice
 - Store $E(x)$ and $E(x^2)$
 - Pre-filter for certain kernel (mipmapping!)
 - Rendering: evaluate $p(d)$ for each fragment


Variance Shadow Maps

[Donnelly and Lauritzen 2006]

shadow Map




variance Shadow Map




Variance Shadow Maps


[Donnelly and Lauritzen 2006]



shadow Map



bilinear PCF




variance Shadow Map

Variance Shadow Maps

[Donnelly and Lauritzen 2006]



SM



PCF 5x5



BiL.PCF 5x5




VSM


Variance Shadow Maps

[Donnelly and Lauritzen 2006]

- $P(d)$ works in many situations
- When σ^2 is large “light bleeding”
 - “Layer” several VSM to alleviate these problems




Revisiting the Shadow Map Test




- $\mathbf{x} \in \mathbb{R}^3$
- $\mathbf{p} \in \mathbb{R}^2$
- \mathbf{x} equals \mathbf{p} just in different spaces
- Shadow function:
 $s(\mathbf{x}) := f(d(\mathbf{x}), z(\mathbf{p}))$
- Binary result:
 - 1 if $d(\mathbf{x}) \leq z(\mathbf{p})$
 - 0 else (shadow)

Shadow Test Function: $s(\mathbf{x})$



- What kind of function is $s(\mathbf{x})$?




- Heaviside Step Function: $H(d-z)$

Convolution Shadow Maps

[Annen et al. 2007]

- Shadow test is a step function
- Idea: transform depth map such that we can write the shadow test as a sum
- Use convolution formula with Fourier expansion



Important Properties of a Fourier Series

- Step function becomes sum of weighted sin() and cos()

$$H(t) \approx c_1 \text{ (plot)} + c_2 \text{ (plot)} + \dots + c_4 \text{ (plot)} + \dots + c_8 \text{ (plot)} + \dots + c_{16} \text{ (plot)}$$

- Series is separable!


$$\sin(d - z) = \boxed{\sin(d)} \boxed{\cos(z)} - \boxed{\cos(d)} \boxed{\sin(z)}$$

- One term depends on shadow map
- One term depends on lookup value
- Pre-filtering possible!

Filtering Example


Original

After filtering




Filtering Example


original



after filtering



Mipmapped CSM recovers fine details (SM: 2048^2)



PCF (hardware)

CSM

CSM – 7x7 Gauss

CSM Blurred Shadows

Filter size: 3x3



SM: 128²

256²

512²

1024²

Filter size: 7x7



SM: 128²


256²

512²

1024²


Issues with a Fourier series

- Ringing suppression
 - Reduce higher frequencies



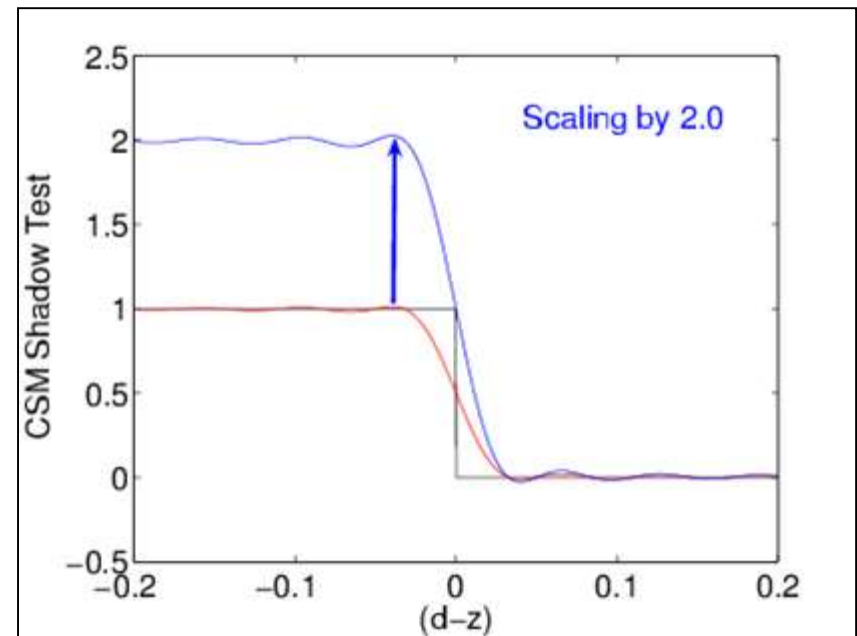
Issues with a Fourier series

- Ringing suppression
 - Reduce higher frequencies
- Steepness of “ramp”
 - Offset (transl. invariance!)
 - Shift shadow test
 - Increases lightness prob.




Issues with a Fourier series

- Ringing suppression
 - Reduce higher frequencies
- Steepness of “ramp”
 - Offset (transl. invariance!)
 - Shift shadow test
 - Increases lightness prob.
 - Scaling
 - Scale shadow test
 - Decreases filtering



Limitations and drawbacks

- Influence of reconstruction order M



$M = 16$

- Memory consumption increases as M grows
- Performance (filtering) decreases as M grows

VSM vs CSM

Variance
Shadow Maps



Convolution
Shadow Maps






Exponential Shadow Maps

[Annen et al. 2008]

- Same general idea as CSM
 - "Linearize" shadow test
- Core idea:
 - Assume $(d(\mathbf{x}) - z(\mathbf{p})) \geq 0$
 - Assume shadow map represents visible front
 - Can use exponential

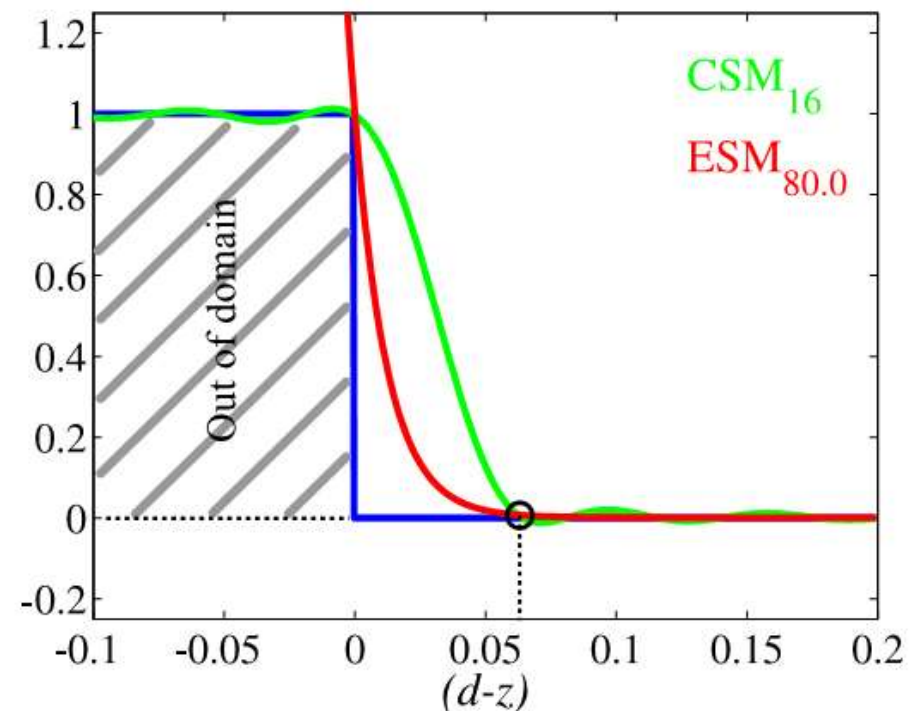

$$f(d(\mathbf{x}), z(\mathbf{p})) \approx e^{-c(d(\mathbf{x}) - z(\mathbf{p}))}$$
$$= e^{-cd(\mathbf{x})} e^{cz(\mathbf{p})}$$



Exponential Shadow Maps


[Annen et al. 2008]

- Same approach as CSM, but uses exponential
- Exponential is separable
- Less memory and faster





Filtering Example










Original e^{cz}



After filtering e^{cz}



Results: Quality Comparison



66 FPS, 21 MB


21 FPS, 170 MB

60 FPS, 42 MB

Faces: 365K, SM resolution: 2Kx2K


Efficient Filtering with Summed Area Tables

- Build summed-area table from the shadow map
 - Can be done efficiently on the GPU
- Summing arbitrary rectangles is $O(1)$!



Efficient Filtering with Summed Area Tables

- Summing arbitrary rectangles



$$average = \frac{LR - UR - LL + UL}{w * h}$$